

Lab7_CartPole_DQN

September 29, 2025

1 Task 1 - Práctica

- Javier Prado - 21486
- Sebastian Juarez - 21471
- Bryan España - 21550 Entrenar a un agente para que equilibre un poste en un carro (**cart-pole**) en movimiento durante el mayor tiempo posible. Se deberá usar **Deep Q-Learning (DQL)** con una **red objetivo** para lograr esto.
- Repositorio de Git: <https://github.com/SebasJuarez/CartPoleDL>

1.1 1. Librerías:

```
[2]: # Imports principales (PyTorch + Gymnasium + NumPy)
import numpy as np
import random
from collections import deque, namedtuple
import matplotlib.pyplot as plt

import torch
import torch.nn as nn
import torch.optim as optim

import gymnasium as gym

# Configuración de dispositivo
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
print("Dispositivo:", device)

# Semillas para reproducibilidad
SEED = 42
np.random.seed(SEED)
random.seed(SEED)
torch.manual_seed(SEED)
if device.type == "cuda":
    torch.cuda.manual_seed_all(SEED)

# Utilidad para gráficos inline
```

```
%matplotlib inline
```

Dispositivo: cuda

1.2 2. Cree el entorno CartPole

```
[3]: # Crear entorno de entrenamiento (sin render para máximo rendimiento)
env = gym.make("CartPole-v1")
env.reset(seed=SEED)

# Opcional: entorno de visualización (crear solo cuando se quiera renderizar
# frames)
# env_vis = gym.make("CartPole-v1", render_mode="rgb_array")
# env_vis.reset(seed=SEED)

print("Observation space:", env.observation_space) # Box(4,)
print("Action space:", env.action_space)          # Discrete(2)
```

Observation space: Box([-4.8 -inf -0.41887903 -inf], [4.8 inf 0.41887903 inf], (4,), float32)
Action space: Discrete(2)

1.3 3. Definan las redes en línea y de destino

```
[4]: class DQN(nn.Module):
    """Red MLP simple para aproximar Q(s,a)."""
    def __init__(self, state_dim, action_dim, hidden_dims=(128, 128)):
        super().__init__()
        layers = []
        last_dim = state_dim
        for h in hidden_dims:
            layers.append(nn.Linear(last_dim, h))
            layers.append(nn.ReLU())
            last_dim = h
        layers.append(nn.Linear(last_dim, action_dim))
        self.net = nn.Sequential(*layers)

    def forward(self, x):
        return self.net(x)

# Inicializar redes (online y target) con misma arquitectura
state_dim = env.observation_space.shape[0]
action_dim = env.action_space.n

q_online = DQN(state_dim, action_dim).to(device)
q_target = DQN(state_dim, action_dim).to(device)

# Copiar pesos al inicio
```

```

q_target.load_state_dict(q_online.state_dict())
q_target.eval()

print(q_online)

```

```

DQN(
  (net): Sequential(
    (0): Linear(in_features=4, out_features=128, bias=True)
    (1): ReLU()
    (2): Linear(in_features=128, out_features=128, bias=True)
    (3): ReLU()
    (4): Linear(in_features=128, out_features=2, bias=True)
  )
)

```

1.4 4. Establecer hiperparámetros

```

[5]: # Hiperparámetros de entrenamiento
GAMMA = 0.99
LR = 1e-3
BATCH_SIZE = 64
BUFFER_SIZE = 100_000
MIN_REPLAY_SIZE = 1_000      # calentamiento del buffer antes de entrenar
TARGET_UPDATE_FREQ = 1000    # cada N pasos (no episodios) se copian pesos a la
    ↪ red objetivo
MAX_EPISODES = 500
MAX_STEPS_PER_EP = 1000

# Exploración epsilon-greedy
EPS_START = 1.0
EPS_END = 0.05
EPS_DECAY_STEPS = 50_000     # pasos para decaer de EPS_START a EPS_END

optimizer = optim.Adam(q_online.parameters(), lr=LR)
loss_fn = nn.MSELoss()

```

1.5 5. Defina la selección de acciones épsilon-greedy

```

[6]: def epsilon_by_step(step):
    """Decaimiento lineal de epsilon con piso en EPS_END."""
    t = min(1.0, step / EPS_DECAY_STEPS)
    return EPS_START + t * (EPS_END - EPS_START)

def select_action(state, epsilon):
    """Selecciona acción según política epsilon-greedy usando q_online."""
    if random.random() < epsilon:
        return env.action_space.sample()

```

```

    else:
        state_v = torch.tensor(state, dtype=torch.float32, device=device).
        ↪unsqueeze(0)
        with torch.no_grad():
            q_vals = q_online(state_v)
            return int(torch.argmax(q_vals, dim=1).item())

```

1.6 6. Defina la reproducción de la experiencia (experience replay)

```

[7]: Transition = namedtuple("Transition", ["state", "action", "reward", "
        ↪next_state", "done"])

class ReplayBuffer:
    def __init__(self, capacity):
        self.buffer = deque(maxlen=capacity)

    def push(self, *args):
        self.buffer.append(Transition(*args))

    def sample(self, batch_size):
        batch = random.sample(self.buffer, batch_size)
        states = torch.tensor([b.state for b in batch], dtype=torch.float32,
        ↪device=device)
        actions = torch.tensor([b.action for b in batch], dtype=torch.int64,
        ↪device=device).unsqueeze(-1)
        rewards = torch.tensor([b.reward for b in batch], dtype=torch.float32,
        ↪device=device).unsqueeze(-1)
        next_states = torch.tensor([b.next_state for b in batch], dtype=torch.
        ↪float32, device=device)
        dones = torch.tensor([b.done for b in batch], dtype=torch.float32,
        ↪device=device).unsqueeze(-1)
        return states, actions, rewards, next_states, dones

    def __len__(self):
        return len(self.buffer)

replay_buffer = ReplayBuffer(BUFFER_SIZE)

```

1.7 7. Ciclo de entrenamiento

```

[ ]: def train_dqn():
    global_step = 0
    episode_rewards = []

    # Calentamiento del buffer
    state, _ = env.reset(seed=SEED)

```

```

for _ in range(MIN_REPLAY_SIZE):
    action = env.action_space.sample()
    next_state, reward, terminated, truncated, _ = env.step(action)
    done = terminated or truncated
    replay_buffer.push(state, action, reward, next_state, done)
    state = next_state if not done else env.reset()[0]

# Entrenamiento principal
for ep in range(1, MAX_EPISODES + 1):
    state, _ = env.reset()
    ep_reward = 0.0

    for step in range(MAX_STEPS_PER_EP):
        epsilon = epsilon_by_step(global_step)
        action = select_action(state, epsilon)

        next_state, reward, terminated, truncated, _ = env.step(action)
        done = terminated or truncated
        replay_buffer.push(state, action, reward, next_state, done)

        # Muestreo de minibatch
        states, actions, rewards, next_states, dones = replay_buffer.
↪sample(BATCH_SIZE)

        with torch.no_grad():
            q_next = q_target(next_states).max(dim=1, keepdim=True)[0]
            q_target_val = rewards + (1 - dones) * GAMMA * q_next

        # Predicción de la red online
        q_pred = q_online(states).gather(1, actions)

        loss = loss_fn(q_pred, q_target_val)
        optimizer.zero_grad()
        loss.backward()
        # Clipping opcional de gradiente para estabilidad
        nn.utils.clip_grad_norm_(q_online.parameters(), max_norm=10.0)
        optimizer.step()

        # Actualización de la red objetivo cada N pasos
        if global_step % TARGET_UPDATE_FREQ == 0:
            q_target.load_state_dict(q_online.state_dict())

    ep_reward += reward
    state = next_state
    global_step += 1

```

```

        if done:
            break

    episode_rewards.append(ep_reward)
    if ep % 10 == 0:
        recent = np.mean(episode_rewards[-10:]) if len(episode_rewards) >= 10
    else np.mean(episode_rewards)
    print(f"Episodio {ep:4d} | Recompensa: {ep_reward:6.1f} | Prom(10): {recent:6.1f} | Epsilon: {epsilon:.3f}")

    return episode_rewards

rewards = train_dqn()

```

C:\Users\javil\AppData\Local\Temp\ipykernel_15232\1074164528.py:12: UserWarning: Creating a tensor from a list of numpy.ndarrays is extremely slow. Please consider converting the list to a single numpy.ndarray with numpy.array() before converting to a tensor. (Triggered internally at C:\actions-runner_work\pytorch\pytorch\pytorch\torch\csrc\utils\tensor_new.cpp:256.)

```

states = torch.tensor([b.state for b in batch], dtype=torch.float32,
device=device)

```

| | | | | | | | |
|----------|-----|-------------|------|-----------|------|----------|-------|
| Episodio | 10 | Recompensa: | 23.0 | Prom(10): | 23.6 | Epsilon: | 0.996 |
| Episodio | 20 | Recompensa: | 24.0 | Prom(10): | 18.3 | Epsilon: | 0.992 |
| Episodio | 30 | Recompensa: | 13.0 | Prom(10): | 17.8 | Epsilon: | 0.989 |
| Episodio | 40 | Recompensa: | 16.0 | Prom(10): | 18.7 | Epsilon: | 0.985 |
| Episodio | 50 | Recompensa: | 8.0 | Prom(10): | 20.8 | Epsilon: | 0.981 |
| Episodio | 60 | Recompensa: | 13.0 | Prom(10): | 14.1 | Epsilon: | 0.978 |
| Episodio | 70 | Recompensa: | 13.0 | Prom(10): | 16.1 | Epsilon: | 0.975 |
| Episodio | 80 | Recompensa: | 13.0 | Prom(10): | 23.0 | Epsilon: | 0.971 |
| Episodio | 90 | Recompensa: | 26.0 | Prom(10): | 21.1 | Epsilon: | 0.967 |
| Episodio | 100 | Recompensa: | 22.0 | Prom(10): | 22.0 | Epsilon: | 0.963 |
| Episodio | 110 | Recompensa: | 23.0 | Prom(10): | 22.4 | Epsilon: | 0.959 |
| Episodio | 120 | Recompensa: | 12.0 | Prom(10): | 20.9 | Epsilon: | 0.955 |
| Episodio | 130 | Recompensa: | 28.0 | Prom(10): | 19.7 | Epsilon: | 0.951 |
| Episodio | 140 | Recompensa: | 11.0 | Prom(10): | 18.9 | Epsilon: | 0.947 |
| Episodio | 150 | Recompensa: | 10.0 | Prom(10): | 19.5 | Epsilon: | 0.944 |
| Episodio | 160 | Recompensa: | 21.0 | Prom(10): | 25.0 | Epsilon: | 0.939 |
| Episodio | 170 | Recompensa: | 14.0 | Prom(10): | 25.1 | Epsilon: | 0.934 |
| Episodio | 180 | Recompensa: | 12.0 | Prom(10): | 23.5 | Epsilon: | 0.930 |
| Episodio | 190 | Recompensa: | 17.0 | Prom(10): | 22.6 | Epsilon: | 0.925 |
| Episodio | 200 | Recompensa: | 10.0 | Prom(10): | 19.8 | Epsilon: | 0.922 |
| Episodio | 210 | Recompensa: | 13.0 | Prom(10): | 20.3 | Epsilon: | 0.918 |
| Episodio | 220 | Recompensa: | 19.0 | Prom(10): | 23.1 | Epsilon: | 0.913 |
| Episodio | 230 | Recompensa: | 17.0 | Prom(10): | 24.4 | Epsilon: | 0.909 |
| Episodio | 240 | Recompensa: | 33.0 | Prom(10): | 25.0 | Epsilon: | 0.904 |
| Episodio | 250 | Recompensa: | 17.0 | Prom(10): | 28.4 | Epsilon: | 0.899 |

| | | | | | | | |
|----------|-----|-------------|------|-----------|------|----------|-------|
| Episodio | 260 | Recompensa: | 18.0 | Prom(10): | 20.3 | Epsilon: | 0.895 |
| Episodio | 270 | Recompensa: | 13.0 | Prom(10): | 29.7 | Epsilon: | 0.889 |
| Episodio | 280 | Recompensa: | 18.0 | Prom(10): | 31.2 | Epsilon: | 0.883 |
| Episodio | 290 | Recompensa: | 21.0 | Prom(10): | 21.6 | Epsilon: | 0.879 |
| Episodio | 300 | Recompensa: | 9.0 | Prom(10): | 34.6 | Epsilon: | 0.872 |
| Episodio | 310 | Recompensa: | 10.0 | Prom(10): | 29.3 | Epsilon: | 0.867 |
| Episodio | 320 | Recompensa: | 44.0 | Prom(10): | 30.2 | Epsilon: | 0.861 |
| Episodio | 330 | Recompensa: | 42.0 | Prom(10): | 27.4 | Epsilon: | 0.856 |
| Episodio | 340 | Recompensa: | 11.0 | Prom(10): | 32.2 | Epsilon: | 0.850 |
| Episodio | 350 | Recompensa: | 46.0 | Prom(10): | 35.5 | Epsilon: | 0.843 |
| Episodio | 360 | Recompensa: | 13.0 | Prom(10): | 26.7 | Epsilon: | 0.838 |
| Episodio | 370 | Recompensa: | 11.0 | Prom(10): | 21.2 | Epsilon: | 0.834 |
| Episodio | 380 | Recompensa: | 24.0 | Prom(10): | 29.9 | Epsilon: | 0.828 |
| Episodio | 390 | Recompensa: | 72.0 | Prom(10): | 36.3 | Epsilon: | 0.821 |
| Episodio | 400 | Recompensa: | 46.0 | Prom(10): | 32.6 | Epsilon: | 0.815 |
| Episodio | 410 | Recompensa: | 16.0 | Prom(10): | 34.2 | Epsilon: | 0.809 |
| Episodio | 420 | Recompensa: | 49.0 | Prom(10): | 39.9 | Epsilon: | 0.801 |
| Episodio | 430 | Recompensa: | 93.0 | Prom(10): | 41.3 | Epsilon: | 0.793 |
| Episodio | 440 | Recompensa: | 24.0 | Prom(10): | 31.4 | Epsilon: | 0.787 |
| Episodio | 450 | Recompensa: | 23.0 | Prom(10): | 26.2 | Epsilon: | 0.782 |
| Episodio | 460 | Recompensa: | 23.0 | Prom(10): | 27.3 | Epsilon: | 0.777 |
| Episodio | 470 | Recompensa: | 14.0 | Prom(10): | 21.7 | Epsilon: | 0.773 |
| Episodio | 480 | Recompensa: | 28.0 | Prom(10): | 31.7 | Epsilon: | 0.767 |
| Episodio | 490 | Recompensa: | 14.0 | Prom(10): | 29.8 | Epsilon: | 0.761 |
| Episodio | 500 | Recompensa: | 29.0 | Prom(10): | 35.2 | Epsilon: | 0.755 |

1.8 8. Representar el entorno

```
[ ]: def demo_render(policy_epsilon=0.0, max_steps=500, seed=SEED):
    """Demostración con render a frames RGB (una corrida)."""
    env_vis = gym.make("CartPole-v1", render_mode="rgb_array")
    state, _ = env_vis.reset(seed=seed)
    total = 0.0
    frames = []

    for _ in range(max_steps):
        eps = policy_epsilon
        action = select_action(state, eps)
        obs, reward, terminated, truncated, _ = env_vis.step(action)
        total += reward
        frames.append(env_vis.render()) # (H, W, 3) numpy array
        state = obs
        if terminated or truncated:
            break

    env_vis.close()
    print(f"Recompensa demo: {total:.1f} (epsilon={policy_epsilon})")
```


```

    return frames

# muestra el primer frame
frames = demo_render(policy_epsilon=1.0, max_steps=50)
plt.imshow(frames[0]); plt.axis('off'); plt.show()

```

Recompensa demo: 18.0 (epsilon=1.0)




1.9 9. Supervisar el entrenamiento

```

[ ]: def plot_rewards(rewards, window=20):
    plt.figure(figsize=(8,4))
    plt.plot(rewards, label='Recompensa por episodio')
    if len(rewards) >= window:
        ma = np.convolve(rewards, np.ones(window)/window, mode='valid')
        plt.plot(range(window-1, window-1+len(ma)), ma, label=f'Media móvil_{window}')
    plt.xlabel('Episodio')
    plt.ylabel('Recompensa')
    plt.title('Progreso de entrenamiento en CartPole (DQN)')
    plt.legend()
    plt.grid(True)
    plt.show()

plot_rewards(rewards)

```

1.10 10. Evalúe el rendimiento

```
[33]: def evaluate_agent(n_episodes=10, seed=SEED):
    scores = []
    for i in range(n_episodes):
        state, _ = env.reset(seed=seed + i)
        total = 0.0
        for _ in range(1000):
            # acción greedy (epsilon ~ 0)
            action = select_action(state, epsilon=0.0)
            state, reward, terminated, truncated, _ = env.step(action)
            total += reward
            if terminated or truncated:
                break
        scores.append(total)
    print(f"Evaluación en {n_episodes} episodios → Promedio: {np.mean(scores):.1f} ± {np.std(scores):.1f}")
    return scores

# Ejemplo tras entrenar:
evaluate_agent(10)
```

Evaluación en 10 episodios → Promedio: 113.9 ± 3.3

```
[33]: [109.0, 119.0, 117.0, 110.0, 115.0, 114.0, 114.0, 115.0, 109.0, 117.0]
```

1.11 11. Fine-Tuning

Algunas ideas para mejorar el desempeño del agente: - **Arquitectura**: cambiar `hidden_dims` (p. ej., (256, 256)) o agregar *layer norm*. - **Optimizador**: probar **AdamW**, modificar LR (1e-4 a 3e-3). - **Exploración**: ajustar `EPS_DECAY_STEPS`, usar *epsilon floor* diferente o *NoisyNets*. - **Objetivo**: usar pérdida Huber en vez de MSE (`nn.SmoothL1Loss`). - **Estabilidad**: **Double DQN** (usar la red online para seleccionar la acción **argmax** en el próximo estado y la target para evaluar su valor). - **Actualización target**: probar *soft update* ($\sim 5e-3$) en lugar de copia dura cada N pasos. - **Batching**: aumentar `BATCH_SIZE` si tienes GPU con memoria suficiente. - **Early stopping/Checkpointing**: guardar los mejores pesos cuando la media móvil supere cierto umbral.

1.11.1 Fine-Tuning en acción

Aplicamos una variante **Dueling Double DQN** con ajustes específicos para CartPole. Las mejoras clave son:

- Arquitectura *dueling* más profunda (256→128) para separar estimaciones de valor y ventaja.
- **Reward shaping** ligero que incentiva mantener el poste vertical y penaliza ángulos extremos.
- **Buffer de experiencia priorizado** ($\alpha = 0.6$) con corrección de sesgo (*scheduling*).
- Optimización con **AdamW** (lr = 1e-3, weight decay = 1e-4) y clipping de gradiente moderado.
- Exploración controlada (`_end = 0.01`, decaimiento en 25k pasos) y entrenamiento extendido a 600 episodios.
- Comparación cuantitativa con el modelo base para verificar la mejora real del desempeño.

```
[34]: # Arquitectura dueling + configuración de hiperparámetros refinados
class DuelingDQN(nn.Module):
    def __init__(self, state_dim, action_dim, hidden_dims=(256, 128)):
        super().__init__()
        layers = []
        last_dim = state_dim
        for h in hidden_dims:
            layers.append(nn.Linear(last_dim, h))
            layers.append(nn.ReLU())
            last_dim = h
        self.feature_extractor = nn.Sequential(*layers)

        self.value_head = nn.Sequential(
            nn.Linear(last_dim, 128),
            nn.ReLU(),
            nn.Linear(128, 1)
        )
        self.advantage_head = nn.Sequential(
            nn.Linear(last_dim, 128),
            nn.ReLU(),
            nn.Linear(128, action_dim)
        )

        for module in self.modules():
            if isinstance(module, nn.Linear):
```

```

        nn.init.xavier_uniform_(module.weight)
        nn.init.constant_(module.bias, 0.0)

    def forward(self, x):
        features = self.feature_extractor(x)
        value = self.value_head(features)
        advantage = self.advantage_head(features)
        advantage_centered = advantage - advantage.mean(dim=1, keepdim=True)
        return value + advantage_centered

q_online_ft = DuelingDQN(state_dim, action_dim).to(device)
q_target_ft = DuelingDQN(state_dim, action_dim).to(device)
q_target_ft.load_state_dict(q_online_ft.state_dict())
q_target_ft.eval()

GAMMA_FT = 0.995
LR_FT = 1e-3
BATCH_SIZE_FT = 128
BUFFER_SIZE_FT = 100_000
MIN_REPLAY_SIZE_FT = 3_000
TARGET_UPDATE_FREQ_FT = 300
EPS_START_FT = 1.0
EPS_END_FT = 0.01
EPS_DECAY_STEPS_FT = 25_000
N_EPISODES_FT = 600

optimizer_ft = optim.AdamW(q_online_ft.parameters(), lr=LR_FT,
    ↪weight_decay=1e-4, eps=1e-7)
loss_fn_ft = nn.SmoothL1Loss(beta=1.0)

print("Red dueling inicializada y lista para fine-tuning (configuración
    ↪refinada).")

```

Red dueling inicializada y lista para fine-tuning (configuración refinada).

```

[35]: # Utilidades de exploración y buffer priorizado para el fine-tuning

def epsilon_by_step_ft(step):
    ratio = min(1.0, step / EPS_DECAY_STEPS_FT)
    return EPS_START_FT + ratio * (EPS_END_FT - EPS_START_FT)

def beta_by_step_ft(step, beta_start=0.4, beta_end=1.0, beta_steps=200_000):
    ratio = min(1.0, step / beta_steps)
    return beta_start + ratio * (beta_end - beta_start)

```

```

def select_action_ft(state, epsilon):
    if random.random() < epsilon:
        return env.action_space.sample()
    state_v = torch.tensor(state, dtype=torch.float32, device=device).
    ↪unsqueeze(0)
    with torch.no_grad():
        q_values = q_online_ft(state_v)
    return int(torch.argmax(q_values, dim=1).item())

class PrioritizedReplayBuffer:
    def __init__(self, capacity, alpha=0.6):
        self.capacity = capacity
        self.alpha = alpha
        self.buffer = deque(maxlen=capacity)
        self.priorities = deque(maxlen=capacity)
        self.eps = 1e-6

    def push(self, *transition, priority=None):
        self.buffer.append(Transition(*transition))
        if priority is None:
            priority = max(self.priorities, default=1.0)
        self.priorities.append(priority)

    def sample(self, batch_size, beta):
        probs = np.array(self.priorities, dtype=np.float32)
        probs = probs ** self.alpha
        probs /= probs.sum()

        indices = np.random.choice(len(self.buffer), batch_size, p=probs)
        samples = [self.buffer[idx] for idx in indices]

        weights = (len(self.buffer) * probs[indices]) ** (-beta)
        weights /= weights.max()

        states = torch.tensor([s.state for s in samples], dtype=torch.float32, ↪
    ↪device=device)
        actions = torch.tensor([s.action for s in samples], dtype=torch.int64, ↪
    ↪device=device).unsqueeze(-1)
        rewards = torch.tensor([s.reward for s in samples], dtype=torch.
    ↪float32, device=device).unsqueeze(-1)
        next_states = torch.tensor([s.next_state for s in samples], dtype=torch.
    ↪float32, device=device)
        dones = torch.tensor([s.done for s in samples], dtype=torch.float32, ↪
    ↪device=device).unsqueeze(-1)

```

```

        weights = torch.tensor(weights, dtype=torch.float32, device=device).
↪unsqueeze(-1)

        return states, actions, rewards, next_states, dones, weights, indices

    def update_priorities(self, indices, priorities):
        for idx, priority in zip(indices, priorities):
            self.priorities[idx] = float(priority + self.eps)

    def __len__(self):
        return len(self.buffer)

def shaped_reward(base_reward, next_state, done):
    if done:
        return base_reward - 1.0
    x, _, theta, _ = next_state
    penalty_angle = 0.5 * (abs(theta) / 0.2)
    penalty_position = 0.1 * (abs(x) / 2.4)
    bonus_center = 0.02 if abs(theta) < 0.05 else 0.0
    return base_reward + bonus_center - penalty_angle - penalty_position

```

[]: *# Entrenamiento del agente dueling double DQN con refinamientos*

```

def train_dueling_double_dqn(max_episodes=N_EPISODES_FT, beta_start=0.4):
    replay_buffer = PrioritizedReplayBuffer(BUFFER_SIZE_FT)
    global_step = 0
    episode_rewards_ft = []
    losses_ft = []

    state, _ = env.reset(seed=SEED + 1234)
    for _ in range(MIN_REPLAY_SIZE_FT):
        action = env.action_space.sample()
        next_state, reward, terminated, truncated, _ = env.step(action)
        done = terminated or truncated
        shaped = shaped_reward(reward, next_state, done)
        replay_buffer.push(state, action, shaped, next_state, done)
        state = next_state if not done else env.reset()[0]

    print("Inicio del entrenamiento dueling double DQN con reward shaping y_
↪buffer priorizado...")

    for ep in range(1, max_episodes + 1):
        state, _ = env.reset()
        ep_reward = 0.0

        for step in range(MAX_STEPS_PER_EP):

```

```

epsilon = epsilon_by_step_ft(global_step)
action = select_action_ft(state, epsilon)

next_state, reward, terminated, truncated, _ = env.step(action)
done = terminated or truncated
shaped = shaped_reward(reward, next_state, done)
replay_buffer.push(state, action, shaped, next_state, done)

if len(replay_buffer) >= BATCH_SIZE_FT:
    beta = beta_by_step_ft(global_step, beta_start=beta_start)
    states, actions, rewards, next_states, dones, weights, indices =
    ↪ replay_buffer.sample(BATCH_SIZE_FT, beta)

    with torch.no_grad():
        next_actions = q_online_ft(next_states).argmax(dim=1,
    ↪ keepdim=True)

        q_next = q_target_ft(next_states).gather(1, next_actions)
        targets = rewards + (1 - dones) * GAMMA_FT * q_next

        current_q = q_online_ft(states).gather(1, actions)
        td_errors = targets - current_q

        loss_elements = torch.nn.functional.smooth_l1_loss(current_q,
    ↪ targets, reduction="none")
        loss = (loss_elements * weights).mean()

        optimizer_ft.zero_grad()
        loss.backward()
        nn.utils.clip_grad_norm_(q_online_ft.parameters(), max_norm=3.0)
        optimizer_ft.step()

        priorities = torch.abs(td_errors).detach().cpu().numpy() + 1e-6
        replay_buffer.update_priorities(indices, priorities)

        losses_ft.append(loss.item())

if global_step % TARGET_UPDATE_FREQ_FT == 0:
    q_target_ft.load_state_dict(q_online_ft.state_dict())

    ep_reward += reward
    state = next_state
    global_step += 1

if done:
    break

episode_rewards_ft.append(ep_reward)

```

```

        if ep % 20 == 0:
            recent_avg = np.mean(episode_rewards_ft[-20:])
            current_eps = epsilon_by_step_ft(global_step)
            print(f"Ep {ep:4d} | Recompensa: {ep_reward:6.1f} | Prom(20):␣
↪{recent_avg:6.1f} | : {current_eps:.3f} | Buffer: {len(replay_buffer):6d}")

        return episode_rewards_ft, losses_ft

fine_tuned_rewards, fine_tuned_losses = train_dueling_double_dqn()

```

Inicio del entrenamiento dueling double DQN con reward shaping y buffer priorizado...

C:\Users\javil\AppData\Local\Temp\ipykernel_15232\1408893917.py:58:
DeprecationWarning: Conversion of an array with ndim > 0 to a scalar is deprecated, and will error in future. Ensure you extract a single element from your array before performing this operation. (Deprecated NumPy 1.25.)
self.priorities[idx] = float(priority + self.eps)

| | | | | | | | | | | | | | |
|----|-----|--|-------------|-------|--|-----------|-------|--|---|-------|--|---------|--------|
| Ep | 20 | | Recompensa: | 24.0 | | Prom(20): | 18.1 | | : | 0.986 | | Buffer: | 3363 |
| Ep | 40 | | Recompensa: | 13.0 | | Prom(20): | 20.2 | | : | 0.970 | | Buffer: | 3767 |
| Ep | 60 | | Recompensa: | 19.0 | | Prom(20): | 26.4 | | : | 0.949 | | Buffer: | 4295 |
| Ep | 80 | | Recompensa: | 15.0 | | Prom(20): | 25.6 | | : | 0.928 | | Buffer: | 4808 |
| Ep | 100 | | Recompensa: | 81.0 | | Prom(20): | 29.2 | | : | 0.905 | | Buffer: | 5392 |
| Ep | 120 | | Recompensa: | 29.0 | | Prom(20): | 25.3 | | : | 0.885 | | Buffer: | 5898 |
| Ep | 140 | | Recompensa: | 14.0 | | Prom(20): | 31.2 | | : | 0.860 | | Buffer: | 6523 |
| Ep | 160 | | Recompensa: | 53.0 | | Prom(20): | 27.4 | | : | 0.839 | | Buffer: | 7072 |
| Ep | 180 | | Recompensa: | 36.0 | | Prom(20): | 27.1 | | : | 0.817 | | Buffer: | 7613 |
| Ep | 200 | | Recompensa: | 53.0 | | Prom(20): | 35.6 | | : | 0.789 | | Buffer: | 8325 |
| Ep | 220 | | Recompensa: | 36.0 | | Prom(20): | 37.9 | | : | 0.759 | | Buffer: | 9082 |
| Ep | 240 | | Recompensa: | 144.0 | | Prom(20): | 40.2 | | : | 0.727 | | Buffer: | 9887 |
| Ep | 260 | | Recompensa: | 19.0 | | Prom(20): | 31.1 | | : | 0.703 | | Buffer: | 10509 |
| Ep | 280 | | Recompensa: | 37.0 | | Prom(20): | 61.7 | | : | 0.654 | | Buffer: | 11743 |
| Ep | 300 | | Recompensa: | 136.0 | | Prom(20): | 70.2 | | : | 0.598 | | Buffer: | 13148 |
| Ep | 320 | | Recompensa: | 29.0 | | Prom(20): | 104.4 | | : | 0.515 | | Buffer: | 15236 |
| Ep | 340 | | Recompensa: | 425.0 | | Prom(20): | 165.6 | | : | 0.384 | | Buffer: | 18547 |
| Ep | 360 | | Recompensa: | 335.0 | | Prom(20): | 397.6 | | : | 0.069 | | Buffer: | 26499 |
| Ep | 380 | | Recompensa: | 428.0 | | Prom(20): | 409.9 | | : | 0.010 | | Buffer: | 34696 |
| Ep | 400 | | Recompensa: | 194.0 | | Prom(20): | 348.6 | | : | 0.010 | | Buffer: | 41667 |
| Ep | 420 | | Recompensa: | 500.0 | | Prom(20): | 282.4 | | : | 0.010 | | Buffer: | 47314 |
| Ep | 440 | | Recompensa: | 500.0 | | Prom(20): | 491.1 | | : | 0.010 | | Buffer: | 57135 |
| Ep | 460 | | Recompensa: | 317.0 | | Prom(20): | 437.9 | | : | 0.010 | | Buffer: | 65894 |
| Ep | 480 | | Recompensa: | 500.0 | | Prom(20): | 377.9 | | : | 0.010 | | Buffer: | 73452 |
| Ep | 500 | | Recompensa: | 500.0 | | Prom(20): | 309.5 | | : | 0.010 | | Buffer: | 79642 |
| Ep | 520 | | Recompensa: | 385.0 | | Prom(20): | 465.2 | | : | 0.010 | | Buffer: | 88946 |
| Ep | 540 | | Recompensa: | 500.0 | | Prom(20): | 374.1 | | : | 0.010 | | Buffer: | 96428 |
| Ep | 560 | | Recompensa: | 146.0 | | Prom(20): | 443.1 | | : | 0.010 | | Buffer: | 100000 |

Ep 580 | Recompensa: 148.0 | Prom(20): 127.0 | : 0.010 | Buffer: 100000
 Ep 600 | Recompensa: 500.0 | Prom(20): 251.7 | : 0.010 | Buffer: 100000

```
[ ]: # Evaluación y comparación contra el modelo base

def evaluate_dueling_agent(n_episodes=10, seed=SEED + 5000):
    q_online_ft.eval()
    scores = []
    for i in range(n_episodes):
        state, _ = env.reset(seed=seed + i)
        total = 0.0
        for _ in range(1000):
            state_v = torch.tensor(state, dtype=torch.float32, device=device).
            ↪unsqueeze(0)
            with torch.no_grad():
                action = torch.argmax(q_online_ft(state_v), dim=1).item()
            state, reward, terminated, truncated, _ = env.step(action)
            total += reward
            if terminated or truncated:
                break
        scores.append(total)
    q_online_ft.train()
    print(f"Evaluación dueling: promedio {np.mean(scores):.1f} ± {np.
    ↪std(scores):.1f}")
    return scores

def compare_models_after_finetune(baseline_scores=None, tuned_scores=None,
    ↪n_eval=10):
    if baseline_scores is None:
        print("Evaluando modelo base original...")
        baseline_scores = evaluate_agent(n_eval)
    if tuned_scores is None:
        print("Evaluando modelo dueling fine-tuned...")
        tuned_scores = evaluate_dueling_agent(n_eval)

    base_avg, base_std = np.mean(baseline_scores), np.std(baseline_scores)
    tuned_avg, tuned_std = np.mean(tuned_scores), np.std(tuned_scores)

    print("\n" + "=" * 45)
    print("COMPARACIÓN BASE vs FINE-TUNED")
    print("=" * 45)
    print(f"Modelo base      : {base_avg:6.1f} ± {base_std:4.1f} (min {np.
    ↪min(baseline_scores):.1f}, max {np.max(baseline_scores):.1f})")
    print(f"Modelo fine-tuned : {tuned_avg:6.1f} ± {tuned_std:4.1f} (min {np.
    ↪min(tuned_scores):.1f}, max {np.max(tuned_scores):.1f})")
```



```

improvement_abs = tuned_avg - base_avg
improvement_pct = (improvement_abs / base_avg) * 100 if base_avg > 0 else 0.
↪0
print(f"\nMejora absoluta   : {improvement_abs:+.1f} puntos")
print(f"Mejora relativa    : {improvement_pct:+.1f}%")

return baseline_scores, tuned_scores

baseline_scores, tuned_scores = compare_models_after_finetune()

```

Evaluando modelo base original...

Evaluación en 10 episodios → Promedio: 113.9 ± 3.3

Evaluando modelo dueling fine-tuned...

Evaluación dueling: promedio 500.0 ± 0.0

```

=====
COMPARACIÓN BASE vs FINE-TUNED
=====

Modelo base          : 113.9 ± 3.3 (min 109.0, max 119.0)
Modelo fine-tuned    : 500.0 ± 0.0 (min 500.0, max 500.0)

Mejora absoluta      : +386.1 puntos
Mejora relativa      : +339.0%

```

[]: *# Visualización del progreso tras el fine-tuning*

```

def plot_finetune_progress(baseline_rewards, fine_tuned_rewards,
↪fine_tuned_losses=None, window=20):
    plt.figure(figsize=(14, 5))

    plt.subplot(1, 2, 1)
    plt.plot(baseline_rewards, label='Base (DQN)', color='tab:blue', alpha=0.6)
    plt.plot(fine_tuned_rewards, label='Fine-Tuned (Dueling Double DQN)',
↪color='tab:orange', alpha=0.7)

    if len(baseline_rewards) >= window:
        base_ma = np.convolve(baseline_rewards, np.ones(window)/window,
↪mode='valid')
        plt.plot(range(window-1, window-1+len(base_ma)), base_ma, color='tab:
↪blue', linewidth=2)
    if len(fine_tuned_rewards) >= window:
        ft_ma = np.convolve(fine_tuned_rewards, np.ones(window)/window,
↪mode='valid')
        plt.plot(range(window-1, window-1+len(ft_ma)), ft_ma, color='tab:
↪orange', linewidth=2)

```

```

plt.xlabel('Episodio')
plt.ylabel('Recompensa')
plt.title('Recompensa por episodio')
plt.legend()
plt.grid(alpha=0.3)

plt.subplot(1, 2, 2)
if fine_tuned_losses:
    plt.plot(fine_tuned_losses, color='tab:red', alpha=0.7)
    if len(fine_tuned_losses) >= window:
        loss_ma = np.convolve(fine_tuned_losses, np.ones(window)/window,
        ↪mode='valid')
        plt.plot(range(window-1, window-1+len(loss_ma)), loss_ma,
        ↪color='tab:red', linewidth=2)
        plt.title('Pérdida (Huber) durante el entrenamiento')
        plt.xlabel('Actualización')
        plt.ylabel('Loss')
        plt.grid(alpha=0.3)
    else:
        plt.axis('off')
        plt.text(0.5, 0.5, 'Ejecuta el entrenamiento para visualizar la
        ↪pérdida',
                ha='center', va='center', fontsize=10)

plt.tight_layout()
plt.show()

```

plot_finetune_progress(rewards, fine_tuned_rewards, fine_tuned_losses)

